

# Building a Highly Scalable, Open Source Twitter Clone

Dan Diephouse ([dan@netzoo.id.com](mailto:dan@netzoo.id.com))

Paul Brown ([prb@mult.ifario.us](mailto:prb@mult.ifario.us))

# Motivation

- ★ Wide (and growing) variety of non-relational databases.

(viz. NoSQL – <http://bit.ly/pLhqQ>, <http://bit.ly/17MmTk>)

- ★ Twitter application model presents interesting challenges of scope and scale.

(viz. “Fixing Twitter” <http://bit.ly/2VmZdz>)

# Storage Metaphors

- ★ Key/Value Store

*Opaque values; fast and simple.*

- ★ Examples:

- ★ Cassandra\* – <http://bit.ly/EdUEt>

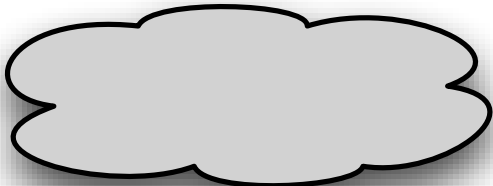
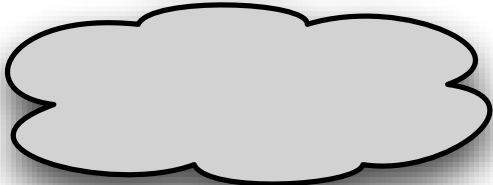
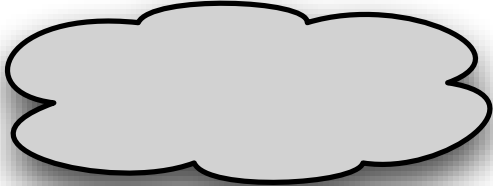
- ★ Dynomite – <http://bit.ly/12AYmf>

- ★ Redis – <http://bit.ly/LBtCh>

- ★ Tokyo Tyrant – <http://bit.ly/oU4uV>

- ★ Voldemort – <http://bit.ly/oU4uV>

# Key/Value

Key	Value
1	
2	
3	

# Storage Metaphors

- ★ Document-Oriented

*Unstructured content; rich queries.*

- ★ Examples:

- ★ CouchDB — <http://bit.ly/JAgUM>

- ★ MongoDB — <http://bit.ly/HDDOV>

- ★ SOLR — <http://bit.ly/q4gyi>

- ★ XML databases...

# Document-Oriented

```
ID="dan-tweet-1",  
TEXT="hello world"
```

```
ID=dan-tweet-2,  
TEXT="Twirp!",  
IN-REPLY-TO="paul-tweet-5"
```

# Storage Metaphors

- ★ Column-Oriented

*Organized in columns; easily scanned.*

- ★ Examples:

- ★ Cassandra★ – <http://bit.ly/EdUEt>

- ★ BigTable – <http://bit.ly/QqMYA>

(available within AppEngine)


- ★ HBase – <http://bit.ly/Zck7F>

- ★ SimpleDB – <http://bit.ly/toh0P>

(Typical library for Java – <http://bit.ly/22kxZ4>)

# Column-Oriented

Name	Date	Tweet Text
Bob	20090506	Eating dinner.
Dan	20090507	Is it Friday yet?
Dan	20090506	Beer me!
Ralph	20090508	My bum itches.



Index	Name
0	Bob
1	Dan
2	Dan
3	Ralph

*Storage*

Index	Date
0	20090506
1	20090507
2	20090506
3	20090508

*Storage*

Index	Tweet Text
0	Eating dinner.
1	Is it Friday yet?
2	Beer me!
3	My bum itches.

*Storage*




# Every Store is Special.

- ★ Lots of different little tweaks to the storage model.
- ★ Widely varying levels of maturity.
- ★ Growing communities.
- ★ Limited (but growing) tooling, libraries, and production adoption.

# Reliability Through Replication

- ★ Consistent hashing to assign keys to partitions.
- ★ Partitions replicated on multiple nodes for redundancy.
- ★ Minimum number of successful reads to consider a write complete.

# Reliability Through Replication



# Web UI

<http://tat1.datapr0n.com:8080>

## bigbird



Isn't flying wonderful? It makes me feel like a bird.

Please sign in:

[\(register\)](#)

username:

admin

password:

••••••••

Login

## bigbird

Find user...

[Home](#) [Sign Out](#)

### What are you doing?

Update

Following:

[user10](#)

[user51](#)

[user11](#)

[user20](#)

[user1](#)

Followers:

Tweets:

[admin](#): Hello World!

07/22/2009 06:07 PM

# Stores

- ★ Tweets

Individual tweets.

- ★ Friends' Timeline

Fixed-length timelines.

- ★ Users

Info and followers.

- ★ Command Queue

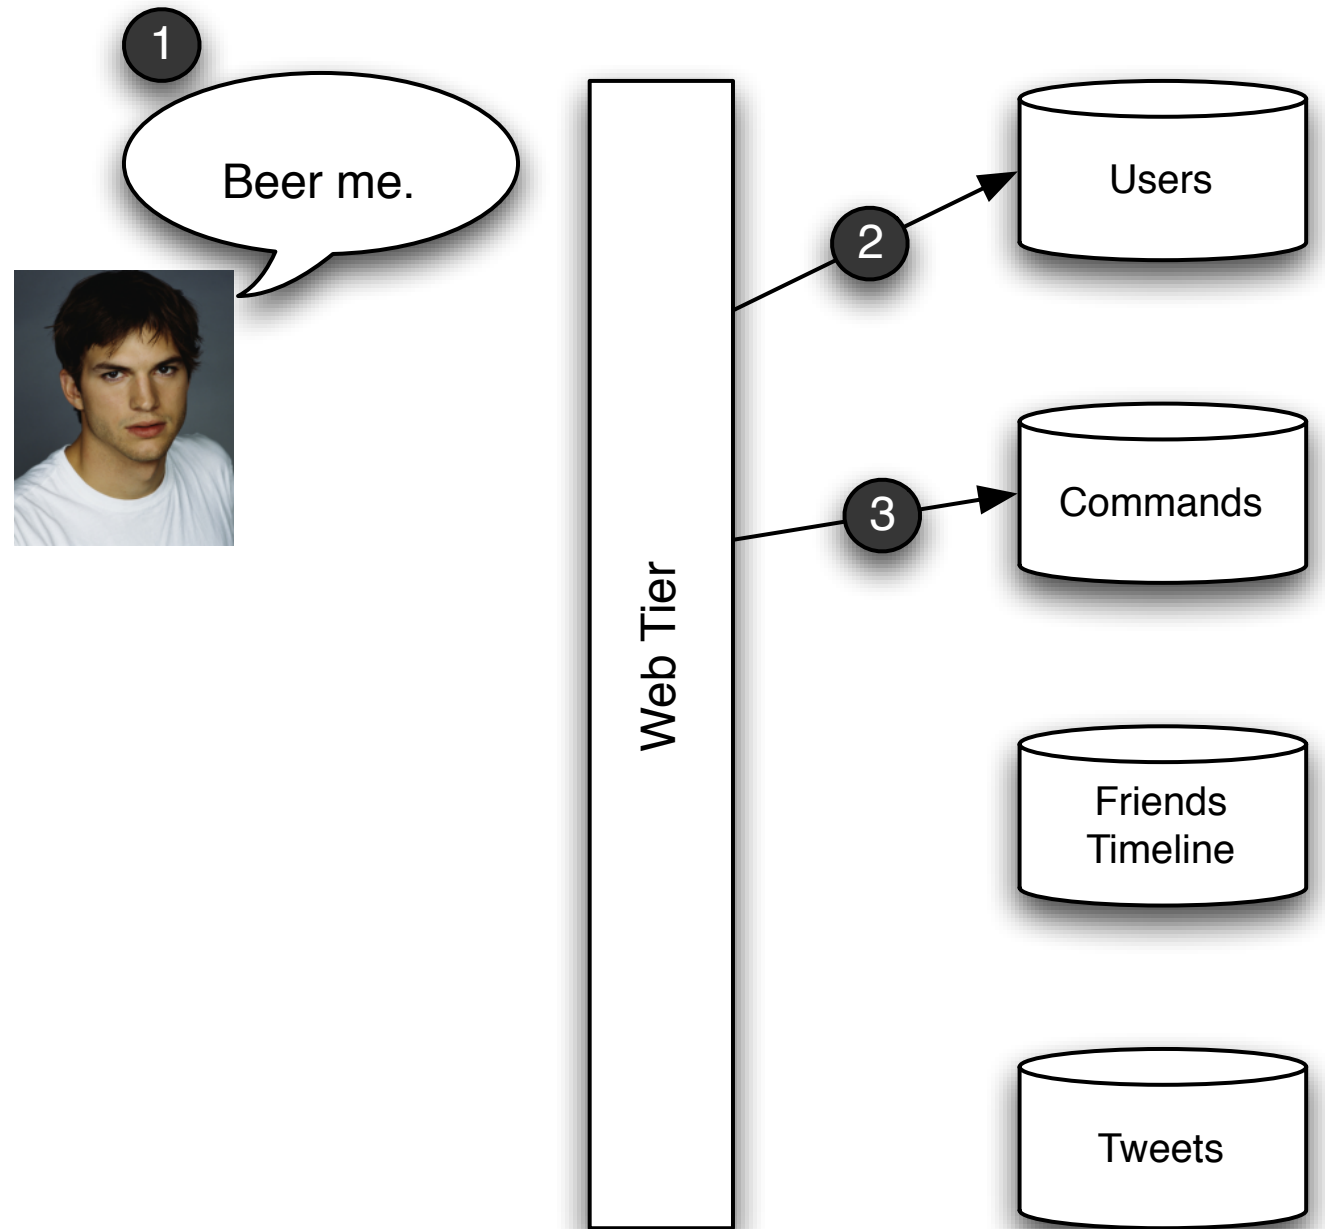
Actions to perform (tweet, follow, etc.).

# Data

- ★ Command (Java serialization)  
Keyed by node name, increasing ID.
- ★ Tweets (Java serialization)  
Keyed by user name, increasing ID.
- ★ FriendsTimeline (Java serialization)  
Keyed by username.  
List of date, tweet ID.
- ★ Users (Java serialization)  
Keyed by username.  
Followers (list), Followed (list), last tweet ID.

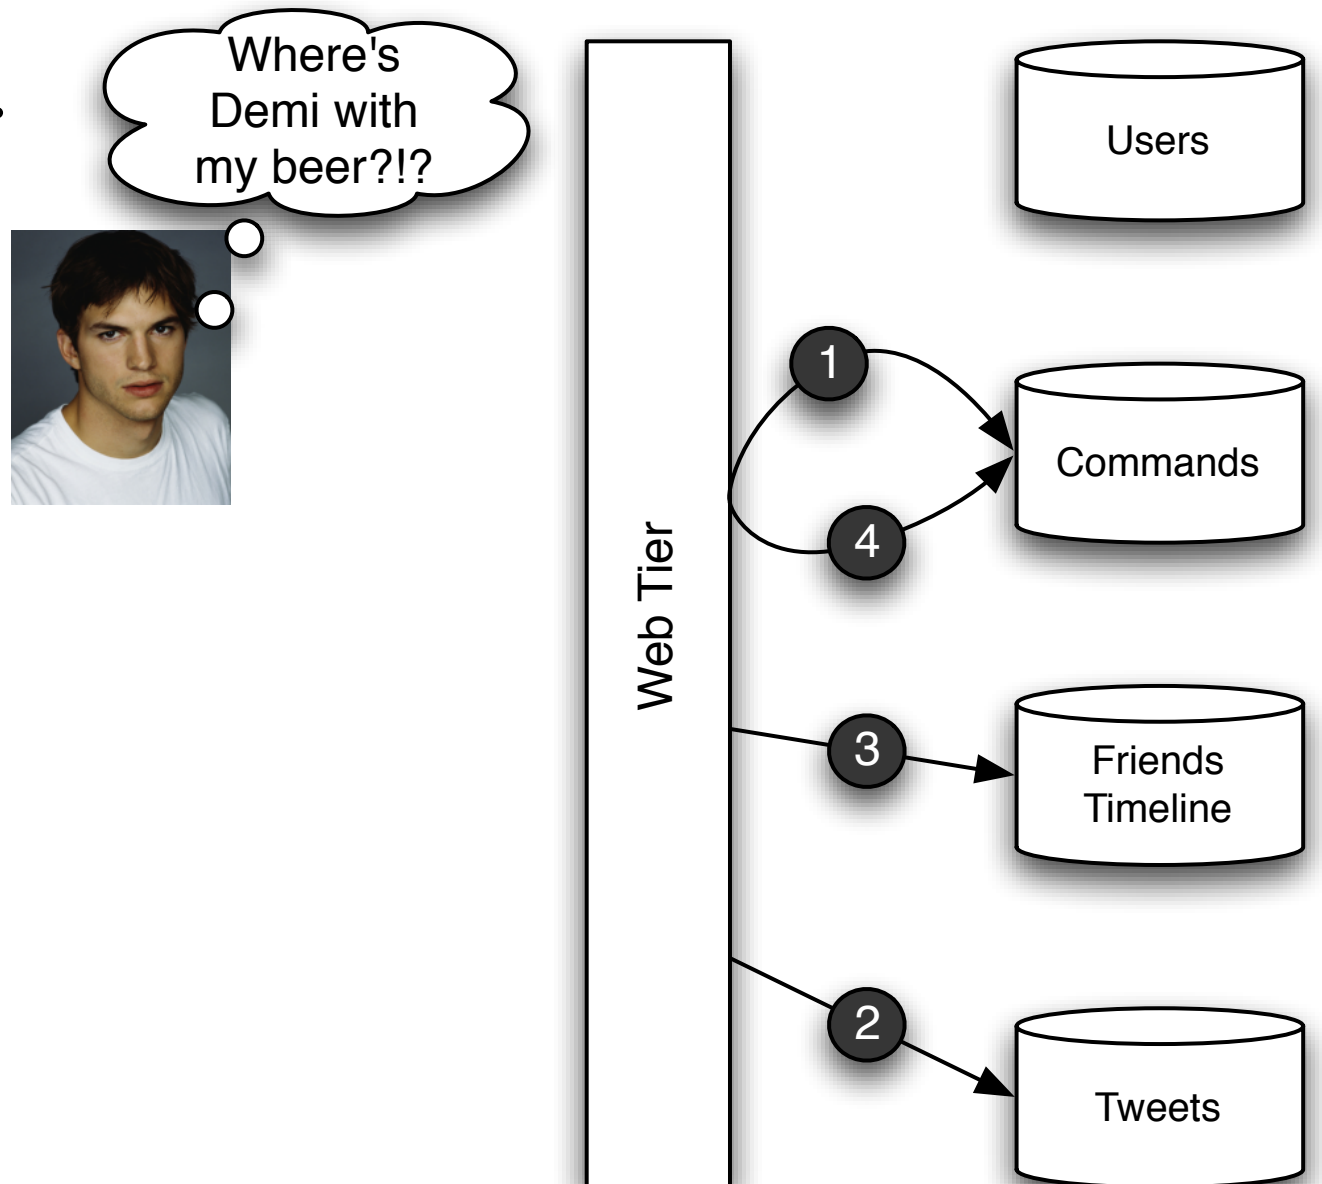
# Life of a Tweet, Part I

1. User tweets.
2. Find next tweet ID for user.
3. Store "tweet for user" command.



# Life of a Tweet, Part II

1. Read next command.
2. Store tweet in user's timeline (Tweets).
3. Store tweet ID in friends' timelines. (Requires \*many\* operations.)
4. DELETE command.





# Some Patterns

- ★ “Sequences” are implemented as race-for-non-collision.
- ★ “Joins” are common keys or keys referenced from values.
- ★ “Transactions” are idempotent operations with DELETE at the end.

# Operations

- ★ Deploy to Amazon EC2
  - ★ 2 nodes for Voldemort
  - ★ 2 nodes for Tomcat
  - ★ 1 node for Cacti
- ★ All “small” instances w/RightScale CentOS 5.2 image.
- ★ Minor inconvenience of “EBS” volume for MySQL for Cacti.

(follow Eric Hammond's tutorial – <http://bit.ly/OK5LZ>)

# Deployment

- ★ Lots of choices for automated rollout (Chef, Capistrano, etc.)
- ★ Took simplest path – Maven build, Ant (scp/ssh and property substitution tasks), and bash scripts.

```
for i in vn1 vn2; do  
    ant -Dnode=${i} setup-v-node  
  
done
```

- ★ Takes ~30 seconds to provision a Tomcat or Voldemort node.

# Dashboarding

- ★ As above, *lots* of choices

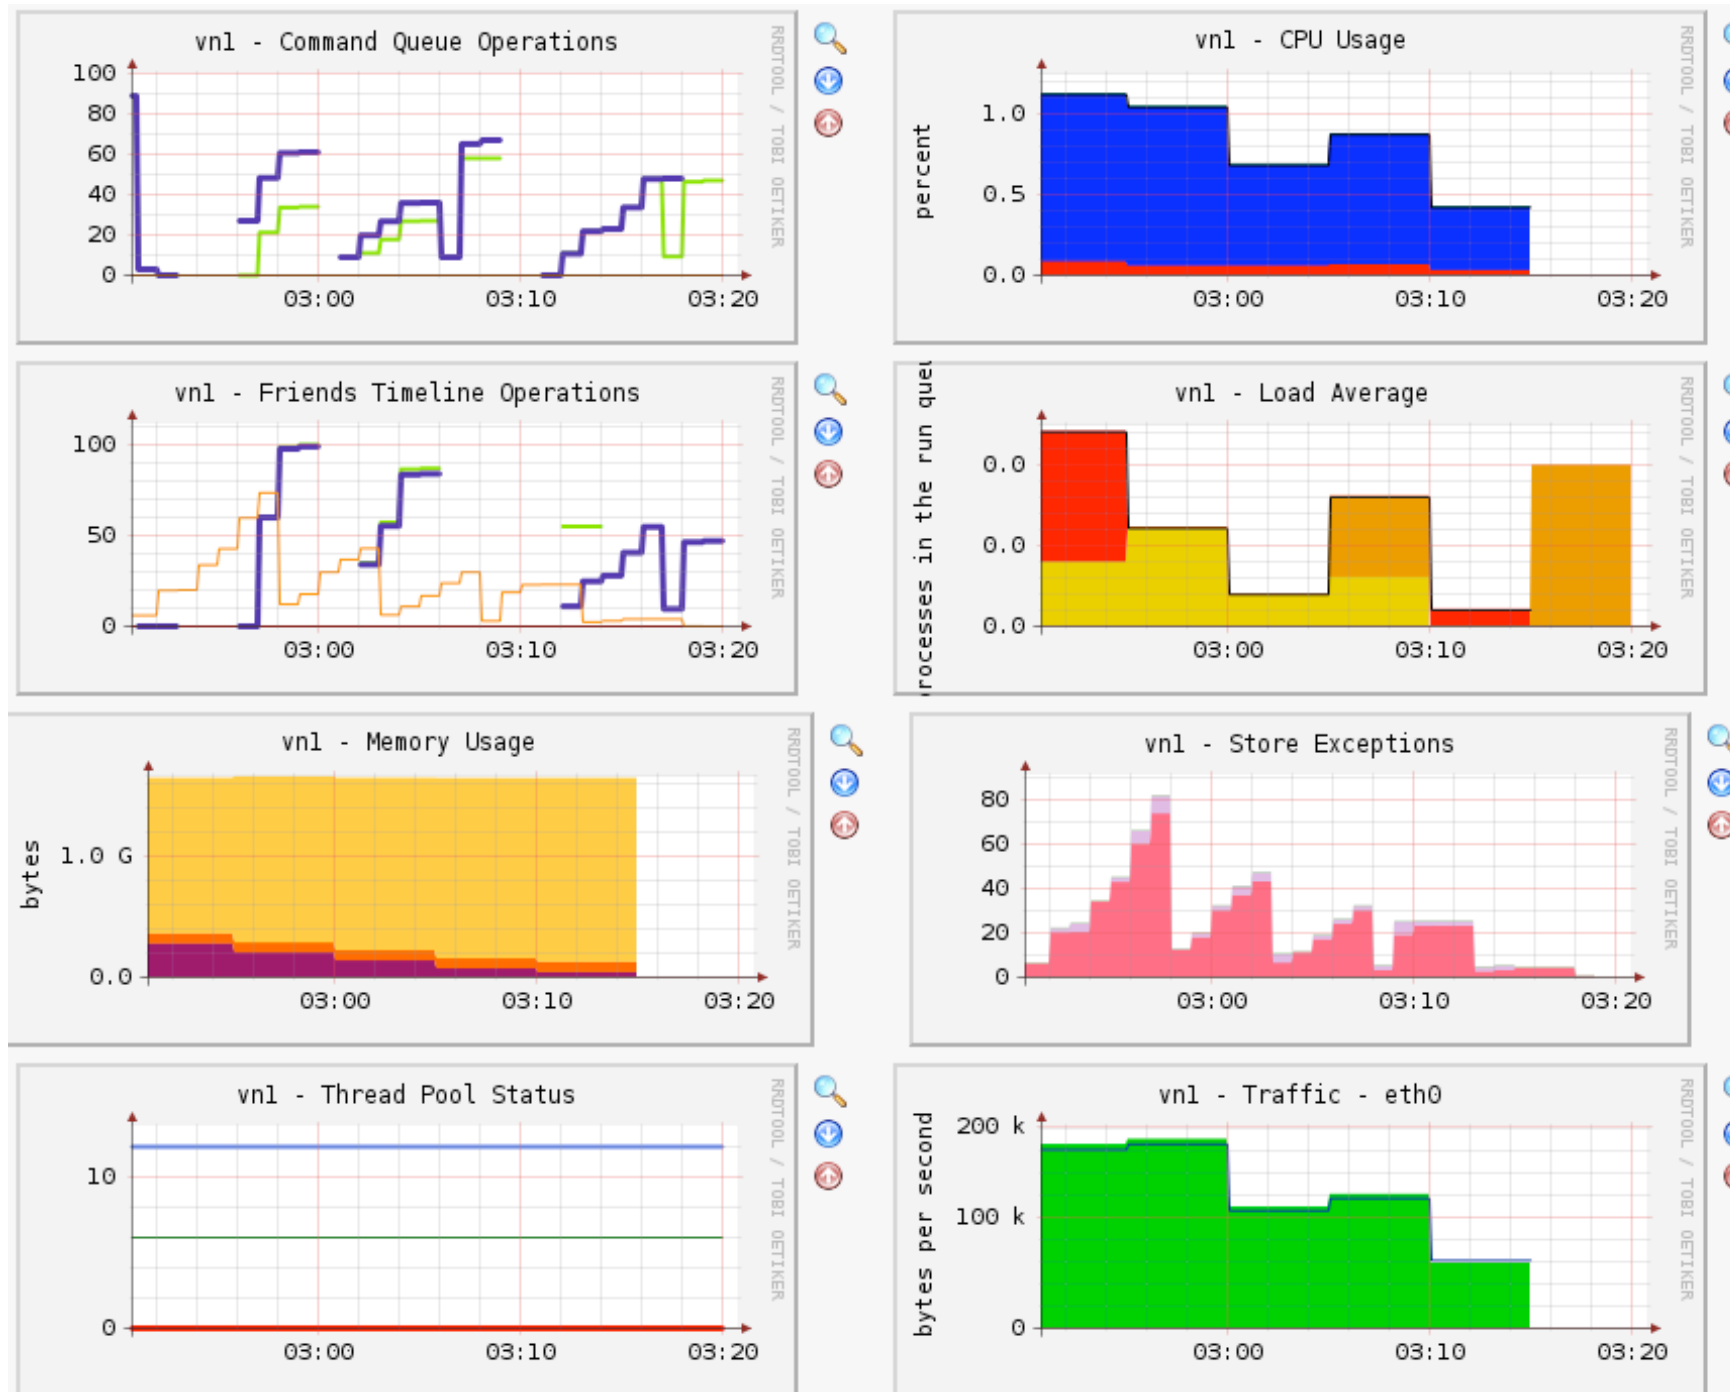
(Cacti – <http://bit.ly/qV4gz>, Graphite – <http://bit.ly/466NAx>, etc.)

- ★ Cacti as simplest choice.

```
yum install -y cacti
```

- ★ Vanilla SNMP on nodes for host data.
- ★ Minimal extensions to Voldemort for stats in Cacti-friendly format.

# Dashboarding



# Performance

- ★ 270 req/sec for getFriendsTimeline against web tier.
- ★ 21 GETs on V stores to pull data.
- ★ 5600 req/sec for V is similar to performance reported at NoSQL meetup (20k req/sec) when adjusted for hardware.
- ★ Cache on the web tier could make this faster...
- ★ Some hassles when hammering individual keys with rapid updates.

# Take Aways

- ★ Linked-list representation deserves some thought (and experiments).

Dynomite + Osmos (<http://bit.ly/BYMdW>)

- ★ Additional use cases (search, rich API, replies, direct messages, etc.) might alter design.
- ★ BigTable/HBase approach deserves another look.
- ★ Source code is available; come and git it.

<http://github.com/prb/bigbird>

<git://github.com/prb/bigbird.git>

# Coordinates

★ Dan Diephouse (@dandiep)

dan@netzoooid.com

http://netzoooid.com

★ Paul Brown (@paulrbrown)

prb@mult.ifario.us

http://mult.ifario.us/a